


# Netflow Aggregation Service

## Background :

The *Netflow Aggregation Service* collects network “flows” (aggregated connection data) to understand how the network is used. The service accepts network flow data points from a set of agents on a large number of instances which monitor outbound connections and report them periodically. The service provides two API's, one for ingesting (write) and the other for reporting (read) netflow data points .



## Endpoints

### Read API

- Path: **/flows**
- Operation: GET
- Query parameters hour (int) - required
- Message/Payload: A JSON array specified below

```
curl "http://localhost:8080/flows?hour=1"
```

### Write API

- Path: **/flows**
- Operation: POST
- Query parameters: None
- Message/Payload: A JSON array specified below

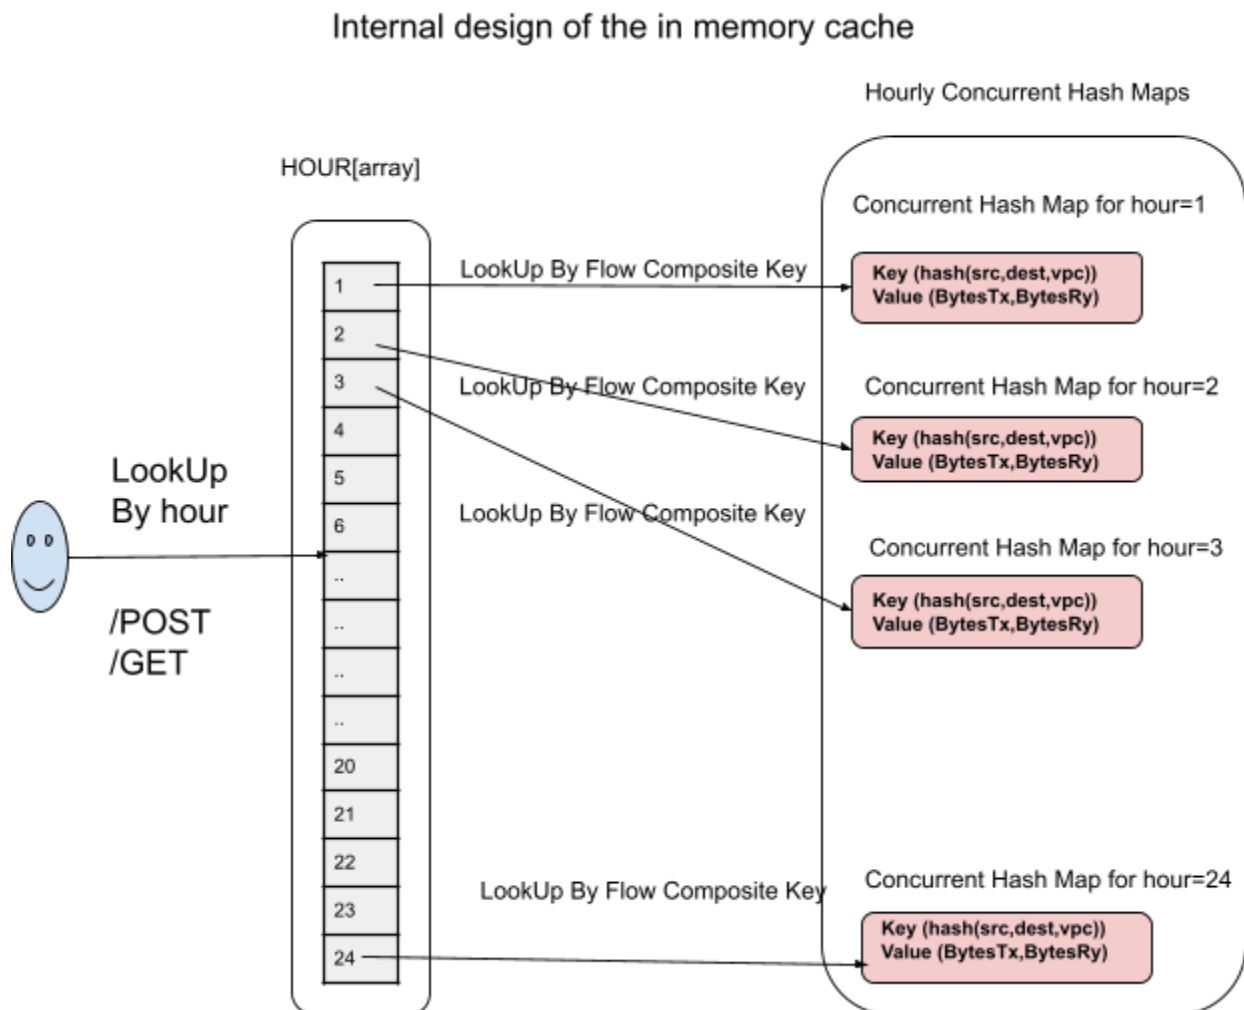
```
curl -X POST "http://localhost:8080/flows" \  
-H 'Content-Type: application/json' \  
-d [{"src_app": "foo", "dest_app": "bar", "vpc_id": "vpc-0", "bytes_tx": 100, "bytes_rx": 500, "hour": 1}]'
```

## Requirements

- 1) Handle frequent heavy write batch loads periodically .
- 2) Write heavy service with frequent updates . Reads are periodic .
- 3) The service does not have to worry about the order in which tuples are processed .
- 4) Aggregation of BytesTx and BytesRy per unique *src\_app,dest\_app,vpc\_id, hour*
- 5) Write/Query performance should not degrade with time and load

## Design

[LOCAL] In Memory Implementation



On POST requests that submit a list of *netflow* data point rows , the service iterates over each of the items parsing the json row into a *NetFlowEntity* . Based on the hour field in the json row a *ConcurrentHashMap* is selected which has the key [hash(src\_app,dest\_app,vpc\_id)] and value is an [Atomic Integer Array] . The value field is intentionally kept as an atomic array so that updates to it are thread safe . This also reduces the space complexity of the cache storage .

On GET requests since the hour param is required , the search space is one single *ConcurrentHashMap* with already aggregated values . Persistence and recovery is implemented in the write ahead log (TODO) and there are no ordering requirements of incoming request data

## Storage

- Outer array of hourly data → Fast lookup based on array indexes
- Hourly Data Cache → Concurrenthashmap per hour i.e. 24 HashMaps

Data Sample : Size of a sample JSON string : 85 Bytes

```
{"src_app": "foo", "dest_app": "bar", "vpc_id": "vpc-0", "bytes_tx": 100, "bytes_rx": 300, "hour": 1}
```

Entry Key :

```
hash(src_app,dest_app,vpc_id)  
Size of Key : 4 bytes
```

Entry Value :

```
Atomic Integer Array  
Size of Value : 12 bytes for BytesTx and 12 Bytes for BytesRy
```

Size per map entry : Key + Value = [ 4B + 12B ] = 16 Bytes

Total Storage Requirements : Unique Keys \* 16 Bytes

## Recovery (TODO)

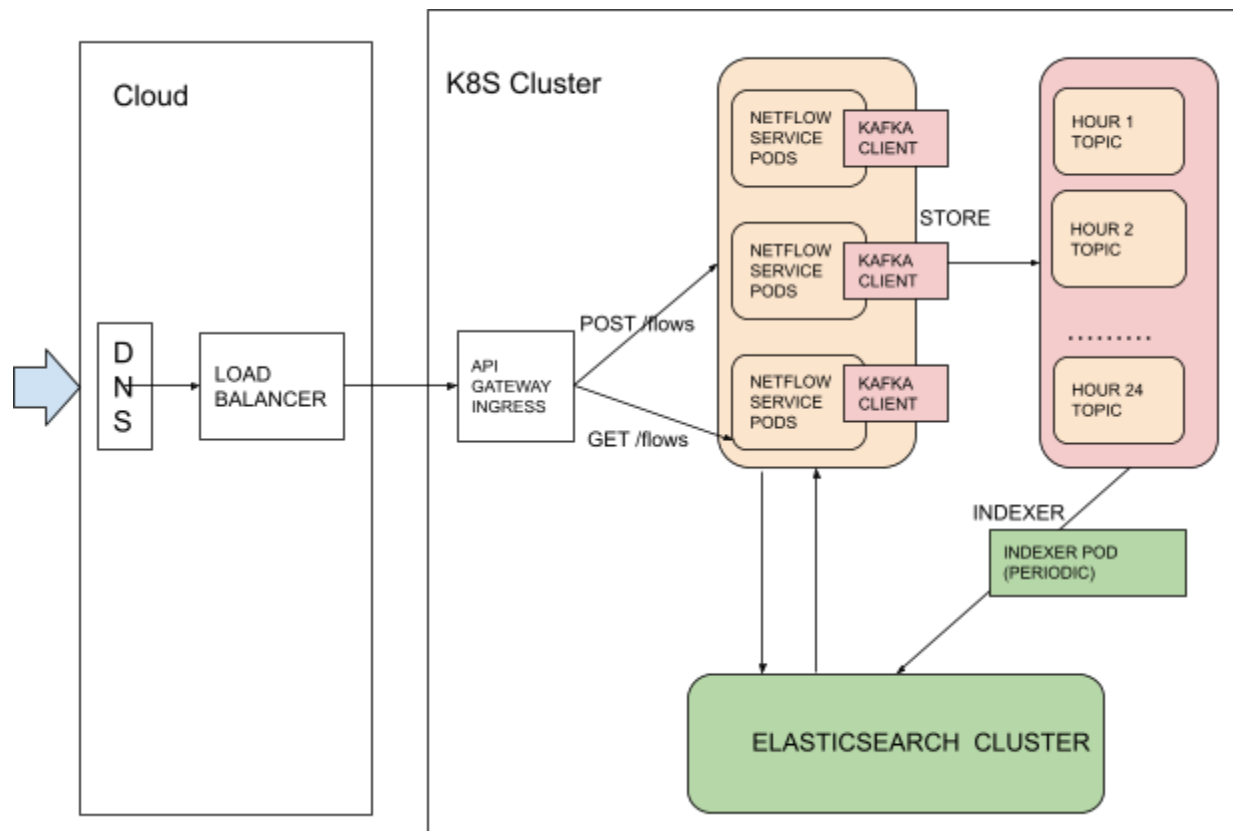
Write ahead log with periodic flushing . Before setting an entry or updating an entry in cache the raw json string can be appended to a list or blocking data structure which can then be flushed at periodic intervals to persist the data to disk . On restarting of the application due to failure this WAL can be replayed to reconstruct the storage structures

## Alternatives (Considered)

Store every tuple without aggregating BytesTx and BytesRy and then on query iterate over the result set and compute/aggregate BytesTx and BytesRy during query processing .

- Advantages : Improves ingestion performance .
- Disadvantages : Degrades query performance .

## [PRODUCTION] In Real World Implementation



- DNS is geo based , hence login from US-WEST1 would hit US-WEST1 Load balancer and same for all regions . This will also enable regional failover using Regional load balancers.
- The Load Balancer is a set of forwarding rules to API proxy.
  - Can be used to terminate HTTPS connections at this layer .
  - Rate Limiting/Request Throttling can be configured at this layer
  - Request Transformation , Request Tracing Headers can be configured here
- Kafka will serve as the source of truth (WRITE AHEAD LOG)
  - Although no ordering is required one TOPIC for writing and reading would suffice
  - Otherwise Kafka can have hourly topics partitioned into tuple composite keys
- Scaling the *service pods* horizontally will enable it to take more load based on traffic .
- Elastic sharded clusters to serve GET requests aggregating data during query phase
- **Redis/Postgres** clusters can also be experimented to see if they can be used instead of Kafka/Elasticsearch